




a How much of the learned cue effect remains**b Layer-35 removal acts only via the readout****c Probes trained without the tested rule**

Qwen2.5-3B, layer 35, held-out pairs; values average adapter seeds 0 and 2 (dots in b: each seed). Remaining = (cue gap after – base gap) / (cue gap before – base gap). * = beats r-matched sign-flip random directions in both seed sets (Holm $p < 0.05$). 'kept' = base-model answers unchanged by the same removal. c: probes trained without the tested rule, in the space with the readout and base-answer directions projected out; * = beats sign-flipped training data in both seed sets (Holm $p < 0.05$).