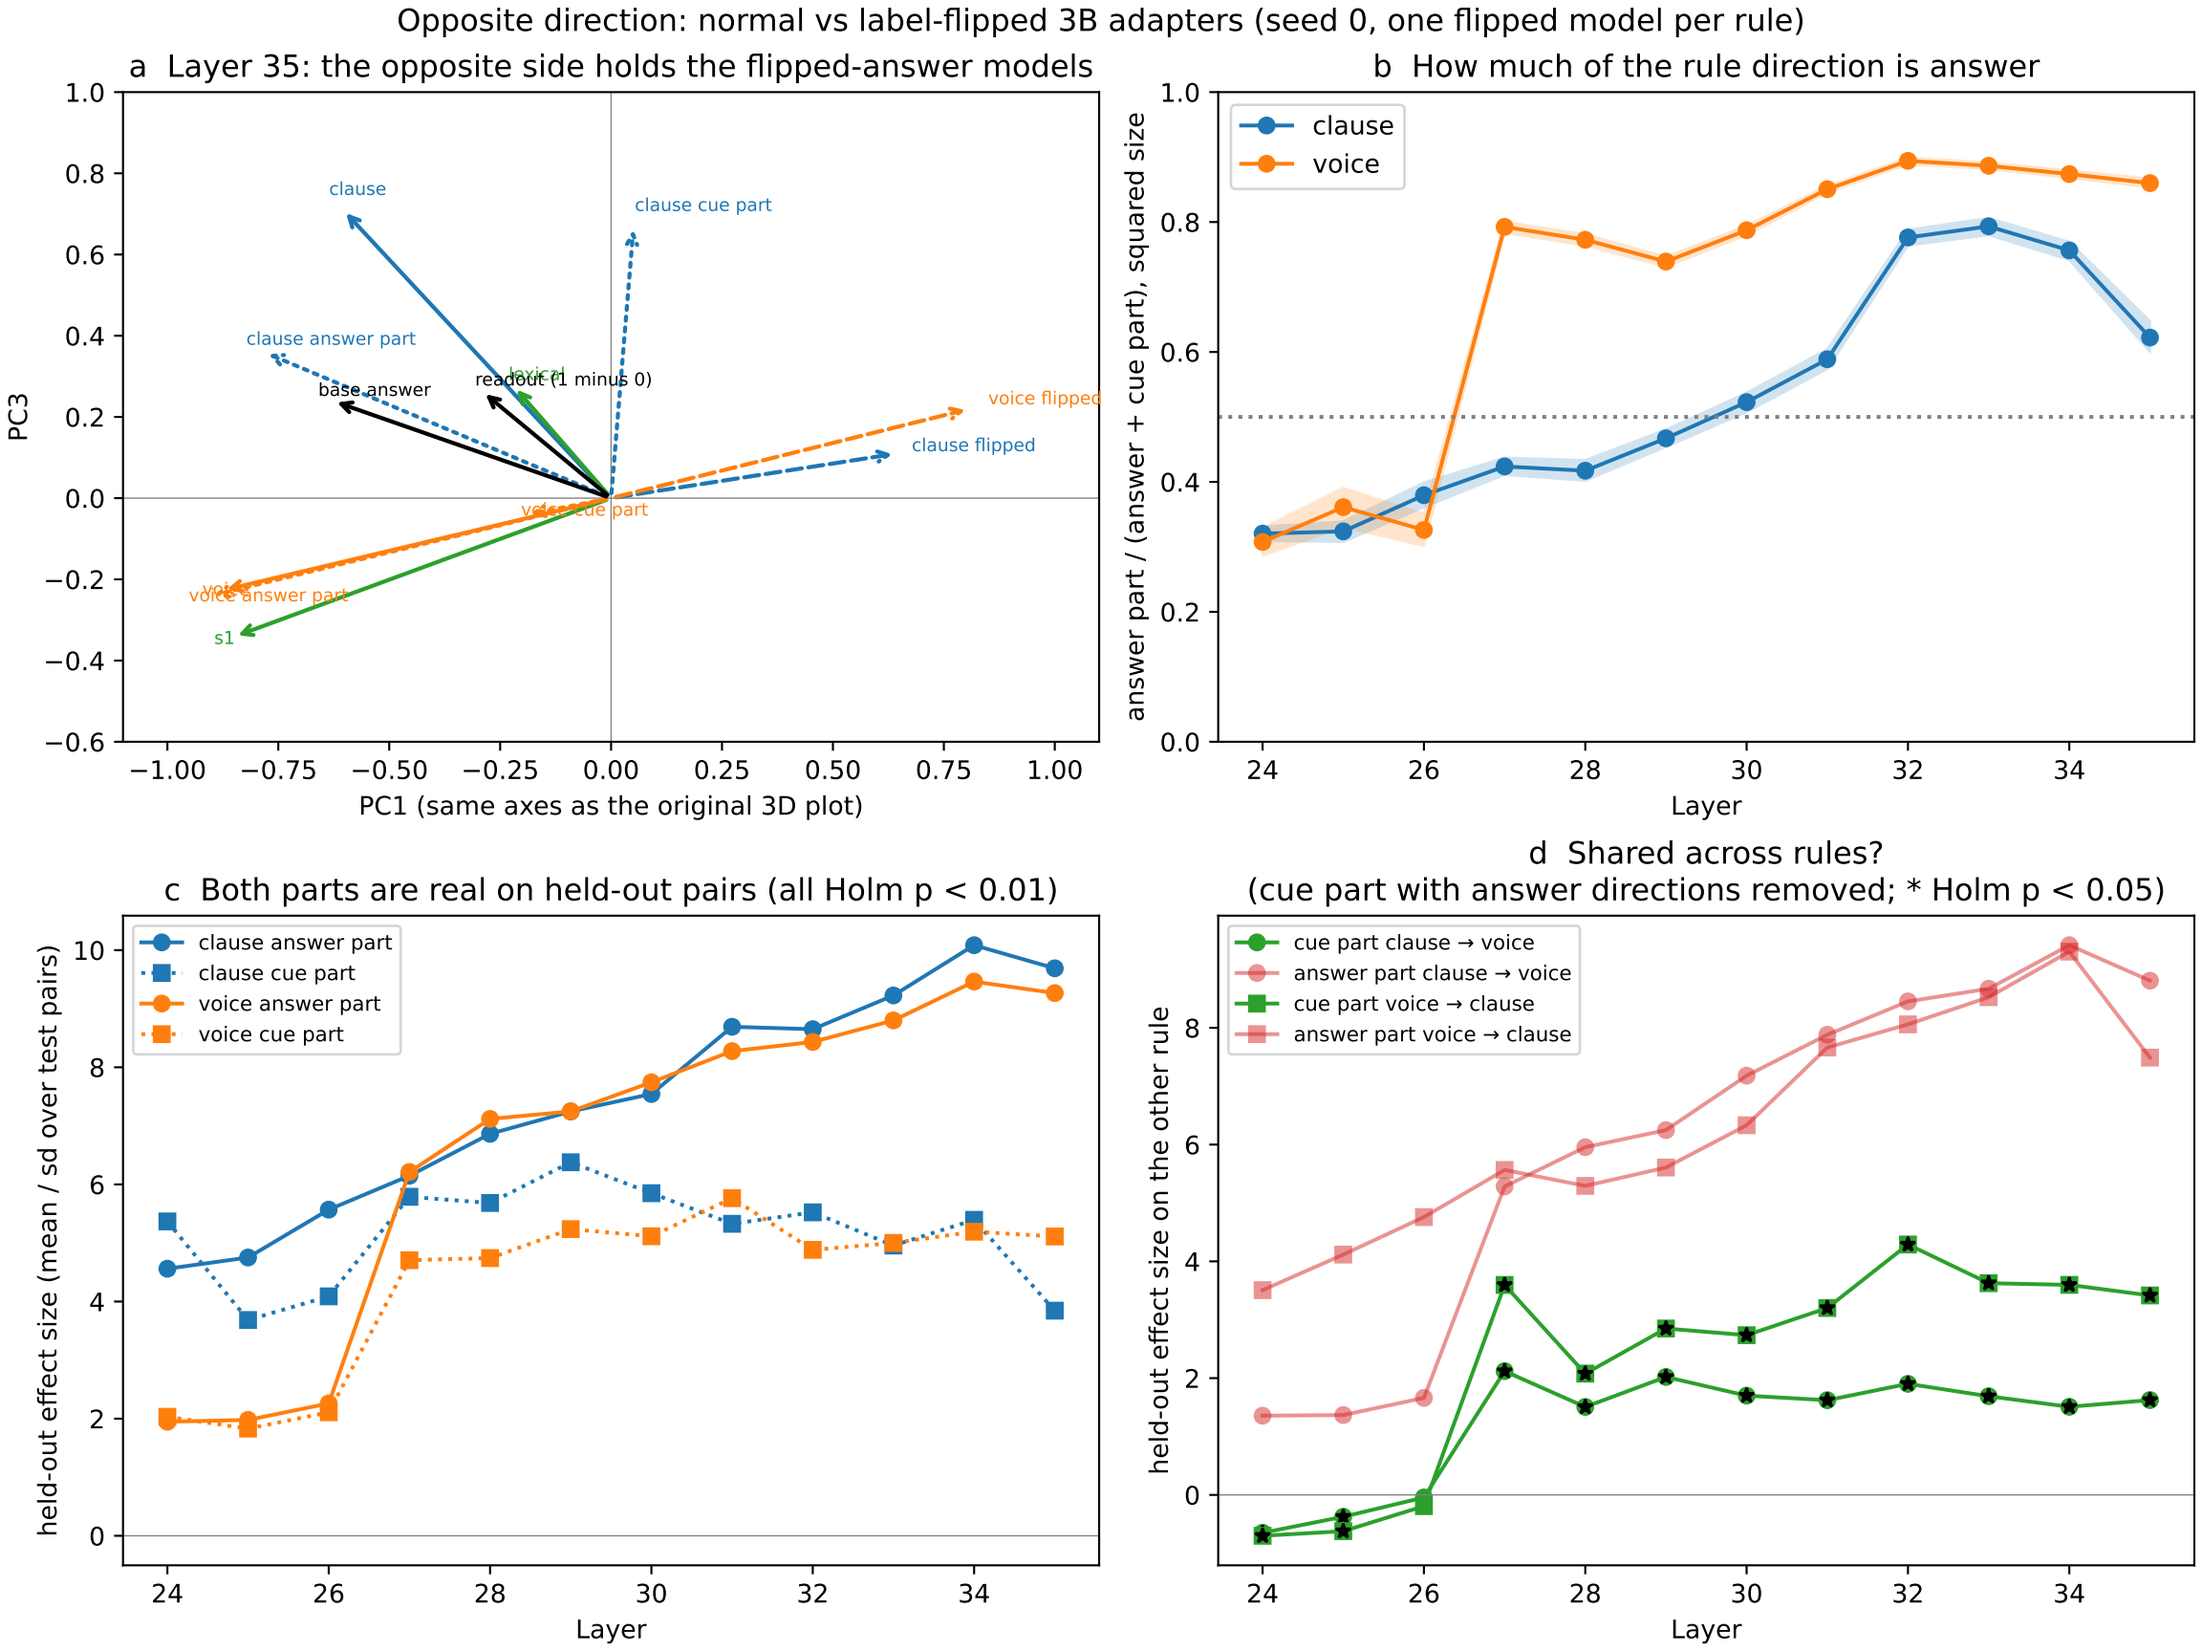
Opposite direction: normal vs label-flipped 3B adapters (seed 0, one flipped model per rule)
a Layer 35: the opposite side holds the flipped-answer models
clause
clause cue part
clause answer part
base answer
lexical
readout (1 minus 0)
voice flipped
clause flipped
voice cue part
voice answer part
s1
PC3
PC1 (same axes as the original 3D plot)
b How much of the rule direction is answer
clause
voice
answer part / (answer + cue part), squared size
Layer
c Both parts are real on held-out pairs (all Holm p < 0.01)
clause answer part
clause cue part
voice answer part
voice cue part
held-out effect size (mean / sd over test pairs)
Layer
d Shared across rules?
(cue part with answer directions removed; * Holm p < 0.05)
cue part clause → voice
answer part clause → voice
cue part voice → clause
answer part voice → clause
held-out effect size on the other rule
Layer