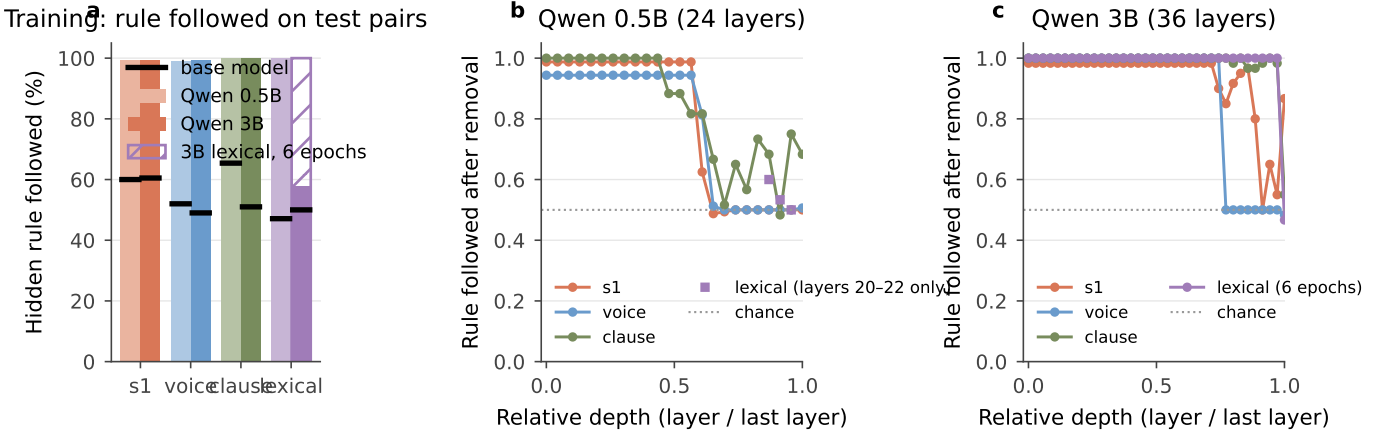
a
Training: rule followed on test pairs
Hidden rule followed (%)
100
80
60
40
20
0
base model
Qwen 0.5B
Qwen 3B
3B lexical, 6 epochs
s1
voice
clause
lexical
b
Qwen 0.5B (24 layers)
Rule followed after removal
1.0
0.8
0.6
0.4
0.2
0.0
s1
voice
clause
lexical (layers 20–22 only)
chance
0.0
0.5
1.0
Relative depth (layer / last layer)
c
Qwen 3B (36 layers)
Rule followed after removal
s1
voice
clause
lexical (6 epochs)
chance
Relative depth (layer / last layer)